

Contents

Introduction	2
Methodology	2
Sorting algorithms to compare against	2
Notable metrics	2
Parameters	3
Data generation	3
Computer Specifications	3
Data Collection Methods	6
Key characteristics	6
Execution pipeline	7
Java implementation	7
Time complexity analysis	9
Best-case scenario	9
Average-case scenario	9
Worst-case scenario	10
Space complexity	11
Comparison with other sorting algorithms	11
Performance results	12
Runtime	12
Comparisons	13
Swaps	14
Copies (excluding swaps)	15
Copies (including swaps)	15
Use Cases & Real-World Examples	16
Conclusion and Limitations	17

Introduction

Timsort is a reliable hybrid sorting algorithm that uses the combination of the Merge Sort and the Binary Insertion Sort. It is an adaptable sorting algorithm, meaning it exploits pre-existing orders sequences (runs) within the input dataset to minimize comparisons and operational overhead, which is different from the $O(n \log_2 n)$ comparison sorts that treat data partially ordered identical and random.

Timsort was designed by Tim Peter in 2002 originally for the Python library, and it was surprisingly successful in handling real life work data and led to global use till these days. Furthermore, Timsort today serves as the default sorting algorithm for object arrays in the Java language (`Arrays.sort`)¹.

This research report evaluates the theoretical mechanics and practical test of our Java execution of Timsort.

Methodology

We will mostly be empirically determining Tim sort's performance through sample trials and data analysis. We will then see whether the data matches with the theoretical time complexity of the algorithm.

We will then compare it against other sorting algorithms to see why it is one of the most implemented sorting algorithm.

Sorting algorithms to compare against

We will compare it against the following sorting algorithms:

- Bubble Sort
- Insertion Sort
- Selection Sort
- Merge Sort
- Timsort
- Java's Sort

The reason why we included Merge Sort is because Timsort builds off of Merge Sort. Moreover, our "control" will be Java's sorting algorithm. Also, do note that our Insertion Sort implementation is Insertion Sort with binary search, that is, Binary Insertion Sort, because, we have to utilize it for Timsort.

Notable metrics

For every trial run, we will keep track of the following metrics:

- runtime
- number of comparisons
- number of swaps

¹The source code for JDK11 for `Arrays.sort` is linked here

- number of copies (excluding swaps)
- number of copies (including swaps)

Swap consists of 3 copies, and since some algorithms don't really use conventional swapping, we believe it's a better metric to compare in place of the amount of swaps. In the code, the `copies` variable is the amount of copies excluding the amount of swaps.

Parameters

Tim sort has a adjustable parameter, which is the minimum run length. For our particular use case, we will use a minimum run size of 32.

Data generation

The most obvious data set we should use is integers, however, that begs the question, what should be the range of the randomly generated integers?

Suppose that we generate random integers in the range between $[1, k]$ and we generate a list of size n . If $k \gg n$, then we would rarely get duplicates, meanwhile, if $k \ll n$, then we would get many duplicates.

Our parameter will then be the ratio $r = \frac{n}{k}$ where a high value of r would mean more duplicates.

In our case, we selected 2 duplication ratios:

- $r = 10$ for data with high duplication
- $r = 0.1$ for data with low duplication

For each duplication ratios, we will use datasets of the following sizes:

- $n = 5$ for small datasets (for data with high duplication, $k = 1$, meaning that the data will just be an array containing just 1)
- $n = 1000$ for medium sized datasets
- $n = 1000000$ for large datasets

Computer Specifications

To make sure that the run time results aren't too much affected by other external factors, we ran the same algorithms on the same computer.

The most important variable for run time in this case is probably the CPU, and the speed and size of the RAM.

In our case, we used a CPU with the following `lscpu` output:

```
Architecture:      x86_64
CPU op-mode(s):    32-bit, 64-bit
Address sizes:      39 bits physical, 48 bits virtual
```

```

Byte Order:                Little Endian
CPU(s):                    6
On-line CPU(s) list:      0-5
Vendor ID:                 GenuineIntel
Model name:               Intel(R) Core(TM) i5-9500 CPU @ 3.00GHz
CPU family:               6
Model:                    158
Thread(s) per core:      1
Core(s) per socket:      6
Socket(s):               1
Stepping:                10
Microcode version:       0xfa
CPU(s) scaling MHz:      20%
CPU max MHz:              4400.0000
CPU min MHz:              800.0000
BogoMIPS:                 6000.00
Flags:                    fpu vme de pse tsc msr pae mce cx8 apic sep mtrr pge mca
                        cmov pat pse36 clflush dts acpi mmx fxsr sse sse2 ss ht
                        tm pbe syscall nx pdpe1gb rdtscp lm constant_tsc art ar
                        ch_perfmon pebs bts rep_good nopl xtopology nonstop_tsc
                        cpuid aperfmperf pni pclmulqdq dtes64 monitor ds_cpl vmx
                        smx est tm2 ssse3 sdbg fma cx16 xtpr pdcm pcid sse4_1 s
                        se4_2 x2apic movbe popcnt tsc_deadline_timer aes xsave a
                        vx f16c rdrand lahf_lm abm 3dnowprefetch cpuid_fault epb
                        pti ssbd ibrs ibpb stibp tpr_shadow flexpriority ept vp
                        id ept_ad fsgsbase tsc_adjust bmi1 avx2 smep bmi2 erms i
                        nvpcid mpx rdseed adx smap clflushopt intel_pt xsaveopt
                        xsavec xgetbv1 xsaves dtherm ida arat pln pts hwp hwp_no
                        tify hwp_act_window hwp_epp vnmi md_clear flush_l1d arch
                        _capabilities

Virtualization features:
Virtualization:           VT-x
Caches (sum of all):
L1d:                      192 KiB (6 instances)
L1i:                      192 KiB (6 instances)
L2:                       1.5 MiB (6 instances)
L3:                       9 MiB (1 instance)
NUMA:
NUMA node(s):             1
NUMA node0 CPU(s):       0-5
Vulnerabilities:
Gather data sampling:     Mitigation; Microcode

```

Ghostwrite:	Not affected
Indirect target selection:	Not affected
Itlb multihit:	KVM: Mitigation: Split huge pages
L1tf:	Mitigation; PTE Inversion; VMX conditional cache flushes , SMT disabled
Mds:	Mitigation; Clear CPU buffers; SMT disabled
Meltdown:	Mitigation; PTI
Mmio stale data:	Mitigation; Clear CPU buffers; SMT disabled
Old microcode:	Not affected
Reg file data sampling:	Not affected
Retbleed:	Mitigation; IBRS
Spec rstack overflow:	Not affected
Spec store bypass:	Mitigation; Speculative Store Bypass disabled via prctl
Spectre v1:	Mitigation; usercopy/swapgs barriers and __user pointer sanitization
Spectre v2:	Mitigation; IBRS; IBPB conditional; STIBP disabled; RSB filling; PBRSE-IBRS Not affected; BHI Not affected
Srbds:	Mitigation; Microcode
Tsa:	Not affected
Tsx async abort:	Mitigation; TSX disabled
Vmscape:	Mitigation; IBPB before exit to userspace

As for the RAM, the output of `sudo dmidecode --type 17` is as follows:

Getting SMBIOS data from sysfs.
SMBIOS 3.1.1 present.

Handle 0x000A, DMI type 17, 40 bytes

Memory Device

```

Array Handle: 0x0009
Error Information Handle: Not Provided
Total Width: 64 bits
Data Width: 64 bits
Size: 8 GiB
Form Factor: DIMM
Set: None
Locator: DIMM1
Bank Locator: Not Specified
Type: DDR4
Type Detail: Synchronous
Speed: 2666 MT/s
Manufacturer: 80CE000080CE
Serial Number: 343C8B13

```

Asset Tag: 02193900
Part Number: M378A1K43DB2-CTD
Rank: 1
Configured Memory Speed: 2666 MT/s
Minimum Voltage: Unknown
Maximum Voltage: Unknown
Configured Voltage: 1.2 V

Handle 0x000B, DMI type 17, 40 bytes
Memory Device

Array Handle: 0x0009
Error Information Handle: Not Provided
Total Width: Unknown
Data Width: Unknown
Size: No Module Installed
Form Factor: Unknown
Set: None
Locator: DIMM2
Bank Locator: Not Specified
Type: Unknown
Type Detail: None

The OS might also be a factor. We used a Linux kernel which has a `uname -a` output as follows:

7.2.3-arch1-3

The test would be decently fair when tested one after the other because the CPU and memory utilization for other processes would be similar.

Data Collection Methods

Because of the randomness of the dataset, and the randomness of OS scheduling, it is necessary to run the code multiple time and record the average of all notable metrics alongside with its standard deviation.

We will run perform each run 100 times to ensure a good sample size.

Key characteristics

- Stability: Timsort is stable, it makes sure that the relative order of equal elements remains unchanged during the sorting process. For sorting objects by multiple fields this is extremely important.
- Adaptability: The way that the algorithm adapts to the data exiting structure is so nicely executed. If the segment of the array is already sorted (ascending or descending), it will capitalize on this order to skip redundant comparisons.

- Comparison based: Timsort can sort items as long as the object as a notion of an order, that is, for a type T , as long as there is a notion that $t_1 > t_2$ for $t_1, t_2 \in T$. This is different from something like Radix Sort or Counting Sort which requires the array to be an array of integers.

Execution pipeline

This is roughly how Timsort gets executed in plain English:

1. Run Identification: The algorithm scans the array to search for subarrays that are already sorted. If a descending run is detected, it will quickly reverse it to prevent the worst case of Insertion Sort from executing. In our implementation, this is done by the `findRun` method, which returns the start of the run, and the length of the run.
2. Run Extension: If the algorithm finds that the natural run is shorter than the `minRun`. Timsort will extend it by pulling in adjacent elements using Binary Insertion Sort until the runs reaches the `minRun` threshold.
3. Merge Stack Management: As runs are identified and measured, they are then pushed onto a runs stack. In our case, in order to make debugging easier, we decided to decouple the runs finding and extension stage, from the invariance preserving stack stage, by storing each run inside of a `discoveredRun` array first, which will then be looped through to be pushed to the actual `runStack`. To help maintain the optimal performance, merging is triggered dynamically in order for the `runStack` to maintain two invariants: $s[n - 2] > s[n - 1] + s[n]$ and $s[n - 1] > s[n]$. Maintaining this invariance has the property that similar sized subarrays will be merged, making it the most optimal $O(n \log_2 n)$ stack collapsing.

Java implementation

```
public Run findRun(T[] arr, int startIndex) {
    int currentIndex = startIndex;
    while ((currentIndex <= (arr.length - 2))
        && (sortMetrics.comparator.compare(arr[currentIndex], arr[currentIndex + 1]) == 0)) {
        currentIndex++;
    }
    if (currentIndex == arr.length - 1) {
        return new Run(startIndex, currentIndex - startIndex + 1, true);
    }
    boolean isAscending;
    isAscending = sortMetrics.comparator.compare(arr[currentIndex + 1], arr[currentIndex]) > 0;
    if (!isAscending) {
        sortMetrics.toggleSign();
    }
}
```

```

while ((currentIndex <= (arr.length - 2))
    && (sortMetrics.comparator.compare(arr[currentIndex + 1], arr[currentIndex]) >= 0)) {
    currentIndex++;
}
if (!isAscending) {
    sortMetrics.toggleSign();
}
return new Run(startIndex, currentIndex - startIndex + 1, isAscending);
}

InsertionSort<T> insertionSort = new InsertionSort<T>(sortMetrics);
while (runPtr <= arr.length - 1) {
    Run run = findRun(arr, runPtr);
    int endIndex;
    // second condition is to check whether the subarray is at the boundary of the array
    if (run.length() < minRun && (run.startIndex() + run.length()) < arr.length) {
        endIndex = Math.min(run.startIndex() + minRun - 1, arr.length - 1);
        if (!run.isAscending()) {
            sortMetrics.toggleSign();
        }
        insertionSort.sortSubarray(arr, run.startIndex(), endIndex);
        if (!run.isAscending()) {
            sortMetrics.toggleSign();
        }
        run = new Run(run.startIndex(), endIndex - run.startIndex() + 1, run.isAscending());
    } else {
        endIndex = runPtr + run.length() - 1;
    }
    if (!run.isAscending()) {
        sortMetrics.reverse(arr, runPtr, endIndex);
        run = new Run(run.startIndex(), run.length(), true);
    }
    runPtr = endIndex + 1;
    discoveredRunPtr++;
    discoveredRun[discoveredRunPtr] = run;
}

for (int i = 0; i <= discoveredRunPtr; i++) {
    runStackPtr++;
    runStack[runStackPtr] = discoveredRun[i];
    mergeCollapse(arr);
}
mergeForceCollapse(arr);

```

Time complexity analysis

We won't fully derive the time complexity for Timsort, and will refer to external authoritative derivations when necessary. Moreover, we will assume that the time complexity of the other algorithms needs no further elaboration.

Best-case scenario

The best-case scenario for Timsort is exactly the same as the best case as Insertion Sort, because, in the case of Timsort, we would find that we only have one run, and even though there are a lot of conditionals afterwards, it wouldn't go into any of them and would just return. In such a case, since `findRun` is $O(n)$, the best case scenario is $O(n)$.

Average-case scenario

For random data, the average case behaves similarly to Merge Sort, where we perform insertion sort once we've reached a sub array of size `minRun`, because, for random data, it is usually the case that we don't have any natural runs that's bigger than `minRun` (unless your duplication ratio is quite high).

A non-rigorous back-of-the-envelope analysis can be done to find the average case running time of Timsort.

Essentially, if our minimum run length is r_m , then we split the array until we reach a sub array of length r_m .

For convenience, let's assume that r_m is a power of 2 and the size of the array n is also a power of 2.

If d is the depth of the merge tree, typically:

$$n = 2^d$$

However, in this case, we factor r_m so that:

$$n = r_m \frac{2^d}{r_m}$$

Meaning that:

$$n = r_m 2^{d - \log_2(r_m)}$$

So, we stop splitting at depth $d_n = d - \log_2(r_m)$

$$d_n = d - \log_2(r_m) = \log_2\left(\frac{n}{r_m}\right)$$

At the bottom, for $\frac{n}{r_m}$ subarrays of size r_m , we need to do Binary Insertion Sort which is of sized r_m , meaning that it takes: Cr_m^2 time where C is a constant factor for some operation. Since we have to do it for $\frac{n}{r_m}$ subarrays, we get that to sort the base of the tree we require:

$$Cr_m^2 \times \frac{n}{r_m} = Cnr_m$$

time.

Merging from that depth level d_n to the top would then take:

$$Cnd_n = Cn \log_2\left(\frac{n}{r_m}\right)$$

because, at every level, we require n operations for merging each subarrays.

As such, in total, we would need:

$$Cn \log_2\left(\frac{n}{r_m}\right) + Cnr_m$$

$$Cn \log_2(n) - Cn \log_2(r_m) + Cnr_m$$

$$Cn \log_2(n) + C(r_m - \log_2(r_m))n$$

time.

So the leading term is $n \log_2(n)$, as such, the algorithm is $O(n \log_2(n))$, as long as r_m is independent of n .

Worst-case scenario

In the case of Insertion Sort, we know that the worst case happens precisely when the array is in descending order, because, it would need to perform $n - 1$ right shifts, and then $n - 2$ right shifts, and then $n - 3$ right shifts... which amounts to $C \frac{n(n-1)}{2}$ (plus binary search, which is $O(\log_2(n))$, which is a lower order term so we ignore it) which is $O(n^2)$.

Timsort isn't naive, in that it has mechanisms to detect descending runs and then reverse them accordingly, preventing the worst case of Insertion Sort from affecting it. This makes analyzing the worst case for Timsort quite difficult. As such I will refer to a paper that analyzed the Worst-case complexity of Timsort by Nicholas Auger, et al. ²

In the abstract, it is stated that for Python's particular implemenation of Timsort, it generally runs in:

²On the Worst-Case Complexity of TimSort by Nicolas Auger, Vincent Jugé, Cyril Nicaud, Carine Pivoteau is linked here

$$O(n + n \log_2 \rho)$$

where ρ is the number of runs.

Anyway, we can even suppose that there is an upper bound that $\rho \leq n$, because the total amount of runs has to be less than the length of the array itself:

$$O(n + n \log_2 \rho) \leq O(n + n \log_2 n)$$

So:

$$O(n + n \log_2 \rho) \leq O(n \log_2 n)$$

Meaning that in the worst case, it is also $O(n \log_2 n)$ for Python's particular implementation of Timsort.

Space complexity

The space complexity of this algorithm is $O(n)$ because we need to initialize a temporary array at least half the size of the array in order to perform the merge operation. In our case, to make implementation easier, we decided to use a temporary array that is the same size as the array, making exactly n , which is still $O(n)$.

Comparison with other sorting algorithms

Algorithm	Best Case	Average Case	Worst Case
Bubble Sort	$O(n^2)$	$O(n^2)$	$O(n^2)$
Selection Sort	$O(n^2)$	$O(n^2)$	$O(n^2)$
Insertion Sort	$O(n)$	$O(n^2)$	$O(n^2)$
Merge Sort	$O(n \log_2 n)$	$O(n \log_2 n)$	$O(n \log_2 n)$
Timsort	$O(n)$	$O(n \log_2 n)$	$O(n \log_2 n)$

It is important to note that Bubble sort, Selection sort can have a best case of $O(n^2)$ for sorted array if we implement early breaking mechanisms such as, if there has been no swaps since we've swept through the array, we break, although, the naive implementation (which is what we implemented) would run in $O(n^2)$ even for an already sorted array

Performance results

Note that the $\pm$ uncertainty range is in terms of the sample standard deviation σ . If we assume that it's normally distributed, we can estimate that 99.7% of the sample falls in between $\mu \pm 3\sigma$, where μ is the sample mean. If we remain conservative and assume that the data is of an arbitrary distribution, Chebyshev's theorem gives us that 99% of the samples must fall in between the range of $\mu \pm 10\sigma$.

Also note that we will sometimes refer to the array of sized 5 as small sized, the array of sized 1000 as medium sized, and the array of sized 1000000 as large sized.

Runtime

Note that the numbers are in seconds.

With $r = 0.1$:

Algorithm	5	1000	1000000
bubbleSort	$1.13 \times 10^{-5} \pm 7.90 \times 10^{-5}$	$2.31 \times 10^{-3} \pm 9.53 \times 10^{-4}$	N/A
insertionSort	$5.35 \times 10^{-6} \pm 7.69 \times 10^{-6}$	$8.43 \times 10^{-4} \pm 7.57 \times 10^{-4}$	N/A
javaSort	$4.09 \times 10^{-4} \pm 4.06 \times 10^{-3}$	$2.28 \times 10^{-4} \pm 8.95 \times 10^{-5}$	$2.68 \times 10^{-1} \pm 2.25 \times 10^{-2}$
mergeSort	$5.06 \times 10^{-6} \pm 8.06 \times 10^{-6}$	$6.32 \times 10^{-4} \pm 1.84 \times 10^{-4}$	$4.53 \times 10^{-1} \pm 1.41 \times 10^{-2}$
selectionSort	$3.90 \times 10^{-6} \pm 6.84 \times 10^{-6}$	$1.24 \times 10^{-3} \pm 8.31 \times 10^{-4}$	N/A
timSort	$1.11 \times 10^{-5} \pm 4.69 \times 10^{-5}$	$2.90 \times 10^{-4} \pm 9.38 \times 10^{-5}$	$4.55 \times 10^{-1} \pm 2.59 \times 10^{-2}$

With $r = 10$:

Algorithm	5	1000	1000000
bubbleSort	$3.29 \times 10^{-6} \pm 7.40 \times 10^{-6}$	$3.09 \times 10^{-3} \pm 1.01 \times 10^{-3}$	N/A
insertionSort	$3.24 \times 10^{-6} \pm 6.32 \times 10^{-6}$	$8.74 \times 10^{-4} \pm 6.65 \times 10^{-4}$	N/A
javaSort	$1.47 \times 10^{-6} \pm 4.62 \times 10^{-6}$	$9.27 \times 10^{-5} \pm 1.41 \times 10^{-5}$	$2.73 \times 10^{-1} \pm 2.41 \times 10^{-2}$
mergeSort	$4.87 \times 10^{-7} \pm 9.35 \times 10^{-8}$	$1.70 \times 10^{-4} \pm 1.24 \times 10^{-5}$	$4.50 \times 10^{-1} \pm 2.07 \times 10^{-2}$
selectionSort	$4.73 \times 10^{-6} \pm 1.13 \times 10^{-5}$	$1.31 \times 10^{-3} \pm 8.29 \times 10^{-4}$	N/A
timSort	$2.45 \times 10^{-6} \pm 4.37 \times 10^{-6}$	$1.69 \times 10^{-4} \pm 1.82 \times 10^{-5}$	$4.39 \times 10^{-1} \pm 1.03 \times 10^{-2}$

The reason why some of the values for the 1000000 column is missing is because, those are $O(n^2)$ sorting algorithms, as such, they wouldn't be able to finish in a reasonable amount of time.

This can easily be seen with some back-of-the-envelope calculation.

When we move from an input size of 1000 to 1000000, we're increasing the input size by a factor of 1000. As such, the amount of time it would take would be scaled by a factor of $1000^2 = 10^6$ or a million.

Since it took $2.31 \times 10^{-3}s$ for Bubble Sort to sort a sized 1000 array, for a sized 10^6 array, we would expect the algorithm to take take:

$$2.31 \times 10^{-3} \times 10^6 = 2.31 \times 10^3 = 2310s$$

It would then take around 38.5 minutes to complete. If we were to run this 100 times, it would then take around 3850 minutes to complete, which is around 64.2 hours. As such, for $O(n^2)$ sorting algorithms, which in our case, include Bubble Sort, Selection Sort, and Insertion Sort; they won't have an entry in the array of sized 1000000 column of any of the other metrics either.

For data with low duplication, Timsort and Merge Sort performs very similarly for large input (10^6 length array), although, it's about 2 times slower than Java's implementation of `Arrays.sort`. This is somewhat expected since for data with low duplication, most of the runs will be of size `minRun`, which makes it behave similarly to Merge Sort.

Where Timsort runs somewhat faster than Merge Sort in the case of data with low duplication with sized 1000, although, it might not be a statistically significant difference given the uncertainty.

As for the data with high duplication, the results doesn't seem significantly different, perhaps because the duplication ratio is too low.

All in all, for our particular dataset, Timsort performs very similarly to Merge Sort, and since Timsort is $O(n \log_2 n)$, it can sort an array of sized 1000000 in a reasonable amount of time compared to Bubble Sort, Selection Sort, and Insertion Sort.

Comparisons

With $r = 0.1$:

Algorithm	5	1000	1000000
bubbleSort	10 ± 0	$4.995 \times 10^5 \pm 0$	N/A
insertionSort	13.47 ± 3.00	$1.906 \times 10^4 \pm 44.92$	N/A
javaSort	N/A	N/A	N/A
mergeSort	7.18 ± 0.85	$8.709 \times 10^3 \pm 16.35$	$1.867 \times 10^7 \pm 555.35$
selectionSort	14 ± 0	$5.005 \times 10^5 \pm 0$	N/A
timSort	16.35 ± 2.96	$1.342 \times 10^4 \pm 67.41$	$2.343 \times 10^7 \pm 1.627 \times 10^3$

With $r = 10$:

Algorithm	5	1000	1000000
bubbleSort	10 ± 0	$4.995 \times 10^5 \pm 0$	N/A
insertionSort	4 ± 0	$1.894 \times 10^4 \pm 71.26$	N/A

Algorithm	5	1000	1000000
javaSort	N/A	N/A	N/A
mergeSort	7 ± 0	$8.701 \times 10^3 \pm 16.12$	$1.867 \times 10^7 \pm 611.02$
selectionSort	14 ± 0	$5.005 \times 10^5 \pm 0$	N/A
timSort	4 ± 0	$1.337 \times 10^4 \pm 75.44$	$2.343 \times 10^7 \pm 2.075 \times 10^3$

Note that for Java's sort, we couldn't count the amount of swaps, nor the amount of comparisons because they don't use our particular implementation of swap and comparison which has innate metrics logging.

The amount of comparisons has a standard deviation that's quite tight, which is good for determining the difference between different algorithms.

For the medium sized array, Timsort has fewer comparisons than Insertion Sort, however, Merge Sort has a significantly lower amount of comparisons. This is probably due to the fact that Timsort requires comparisons for merging similar to Merge Sort, but also comparisons to determine the amount of runs, whether it's decreasing, and because we use Insertion Sort for smaller sized array.

Swaps

With $r = 0.1$:

Algorithm	5	1000	1000000
bubbleSort	4.92 ± 1.98	$2.492 \times 10^5 \pm 5.133 \times 10^3$	N/A
insertionSort	0 ± 0	0 ± 0	N/A
javaSort	N/A	N/A	N/A
mergeSort	0 ± 0	0 ± 0	0 ± 0
selectionSort	2.97 ± 0.66	992.43 ± 2.36	N/A
timSort	1.14 ± 1.00	246.72 ± 42.51	$2.500 \times 10^5 \pm 1.366 \times 10^3$

With $r = 10$:

Algorithm	5	1000	1000000
bubbleSort	0 ± 0	$2.467 \times 10^5 \pm 4.822 \times 10^3$	N/A
insertionSort	0 ± 0	0 ± 0	N/A
javaSort	N/A	N/A	N/A
mergeSort	0 ± 0	0 ± 0	0 ± 0
selectionSort	3 ± 0	993.42 ± 2.21	N/A
timSort	0 ± 0	247.52 ± 44.52	$2.499 \times 10^5 \pm 1.244 \times 10^3$

Bubble Sort and Selection Sort requires swaps, however, Insertion Sort and Merge Sort doesn't do a conventional swap, as such, this isn't quite a fair metric of comparison.

Copies (excluding swaps)

With $r = 0.1$:

Algorithm	5	1000	1000000
bubbleSort	0 ± 0	0 ± 0	N/A
insertionSort	11.04 ± 3.38	$2.517 \times 10^5 \pm 4.861 \times 10^3$	N/A
javaSort	N/A	N/A	N/A
mergeSort	24 ± 0	$1.995 \times 10^4 \pm 0$	$3.990 \times 10^7 \pm 0$
selectionSort	0 ± 0	0 ± 0	N/A
timSort	8.03 ± 2.91	$1.940 \times 10^4 \pm 173.41$	$3.941 \times 10^7 \pm 5.546 \times 10^3$

With $r = 10$:

Algorithm	5	1000	1000000
bubbleSort	0 ± 0	0 ± 0	N/A
insertionSort	0 ± 0	$2.500 \times 10^5 \pm 4.928 \times 10^3$	N/A
javaSort	N/A	N/A	N/A
mergeSort	24 ± 0	$1.995 \times 10^4 \pm 0$	$3.990 \times 10^7 \pm 0$
selectionSort	0 ± 0	0 ± 0	N/A
timSort	0 ± 0	$1.929 \times 10^4 \pm 162.64$	$3.941 \times 10^7 \pm 5.402 \times 10^3$

The copies excluding swaps gets counted when we do things such as right shift operators and merging. Nonetheless, it still isn't a fair metric of comparison because, Bubble Sort and Selection Sort only uses swaps, which doesn't get counted into the copies excluding swaps.

Copies (including swaps)

With $r = 0.1$:

Algorithm	5	1000	1000000
bubbleSort	14.76 ± 5.95	$7.477 \times 10^5 \pm 1.540 \times 10^4$	N/A
insertionSort	11.04 ± 3.38	$2.517 \times 10^5 \pm 4.861 \times 10^3$	N/A
javaSort	N/A	N/A	N/A
mergeSort	24 ± 0	$1.995 \times 10^4 \pm 0$	$3.990 \times 10^7 \pm 0$
selectionSort	8.91 ± 1.98	$2.977 \times 10^3 \pm 7.07$	N/A

Algorithm	5	1000	1000000
timSort	11.45 ± 4.47	$2.014 \times 10^4 \pm 224.68$	$4.016 \times 10^7 \pm 6.637 \times 10^3$

With $r = 10$:

Algorithm	5	1000	1000000
bubbleSort	0 ± 0	$7.402 \times 10^5 \pm 1.447 \times 10^4$	N/A
insertionSort	0 ± 0	$2.500 \times 10^5 \pm 4.928 \times 10^3$	N/A
javaSort	N/A	N/A	N/A
mergeSort	24 ± 0	$1.995 \times 10^4 \pm 0$	$3.990 \times 10^7 \pm 0$
selectionSort	9 ± 0	$2.980 \times 10^3 \pm 6.64$	N/A
timSort	0 ± 0	$2.004 \times 10^4 \pm 205.08$	$4.016 \times 10^7 \pm 6.658 \times 10^3$

The copies which include swaps includes swaps into the total amount of copies, where one swap amounts to 3 copies operations.

For arrays with a high low duplication factor, Bubble Sort, and Timsort doesn't do any copy and exits for small arrays. The array consists of just 1, meaning that it is already sorted. Merge Sort on the other hand is naive, and still splits the array, even though it's already sorted. Similarly for Selection Sort, it will still methodically search for the minimum and try to swap it, because it's naive.

For bigger arrays, for both arrays with low and high duplication factor, the amount of copies between Merge Sort and Timsort isn't significant. We believe that Timsort would perform a lot better for data more duplicates, and larger runs (natural runs larger than 32).

Use Cases & Real-World Examples

Since Timsort is a hybrid architecture makes it highly advantageous for specific software, although its memory requirements limit its universal application.

Optimal Scenarios (Appropriate Use Cases):

- Real world Data Processing: in this case data often contains pre-existing sequential patterns (partially sorted data); Timsort capitalizes on these natural runs to approach its best case which is the $O(n)$.
- Stability Critical Tasks: Since Timsort is a stable sorting algorithm (assuming that merging is implemented in a stable fashion, which it can be), it is ideal for sorting complex objects in arrays by multiple fields sequentially without disrupting the original relative order of equal elements.

- High-Level Language Core Libraries: Timsort is well optimized for the standard library implementations handling drivers, with unpredictable user datasets, as evidenced by its integration into Python's `list.sort` method (pre-Python 3.11) , Java's `Arrays.sort` ³.

Unoptimal Scenarios (Inappropriate Use Cases):

- Limited Ram: Timsort require auxiliary memory proportional to the data size, it is inappropriate for embedded system that has only the $O(1)$ space requirement.
- Integer Arrays: For Integer arrays, sorting can be done a lot faster using an algorithm like Radix Sort. We choose an array of Integers for testing due to its simplicity.

Conclusion and Limitations

Timsort is seen as the gold standard in terms of general sorting algorithm. Although this demonstration haven't fully demonstrated Timsort's capability for real world data, we've shown that for random data, it performs just as good as Merge sort, which has an $O(n \log_2 n)$ run time, making it feasible for arrays of very large sizes.

In practice, Timsort also has many more optimizations such as galloping mode for merging, and more advanced techniques which we haven't implemented which may make it a lot faster. In the future, it would be a lot more insightful to use datasets that more resemble real world data, and implement the different optimizations that is used in applications such as in Java's `Arrays.sort`.

³Python moved away from using Timsort to using Power Sort starting from Python version 3.11. More information can be found [here](#).